

Introduction to Algorithms

L6. Tree

Instructor : Kilho Lee

Course Overview

- Algorithmic Analysis
- Divide and Conquer
- Randomized Algorithms
- **Tree Algorithms**
- Graph Algorithms
- Dynamic Programming
- Greedy Algorithms
- Advanced Algorithms

Today's Outline

- Tree Algorithms
 - Tree basics, Binary search tree, Red-black tree
 - Reading: CLRS 12.1, 12.2, 12.3 and 13


Some data structures

for storing objects like **5** (aka, **nodes** with **keys**)

- (Sorted) arrays:



- (Sorted) linked lists:



- Some basic operations:
 - **INSERT, DELETE, SEARCH**

Sorted Arrays

1	2	3	4	5	7	8
---	---	---	---	---	---	---

- $O(n)$ INSERT/DELETE:

1	2	3	4	4.5	7	8
---	---	---	---	-----	---	---


- $O(\log(n))$ SEARCH:

1	2	3	4	5	7	8
---	---	---	---	---	---	---




eg, Binary search to see if 3 is in A.

Sorted linked lists




- $O(1)$ INSERT/DELETE:

- (assuming we have a pointer to the location of the insert/delete)









- $O(n)$ SEARCH:




Motivation for Binary Search Trees

TODAY!

	Sorted Arrays	Linked Lists	Binary Search Trees*
Search	$O(\log(n))$ 	$O(n)$ 	$O(\log(n))$ 
Insert/Delete	$O(n)$ 	$O(1)$ 	$O(\log(n))$ 


Binary Search Trees

Recall the data structure



- Graph
 - Graph $G = (V, E)$ is a pair of the vertex set V and the edge set of E

Recall the data structure




- Tree

- Let $G = (V, E)$ be an undirected graph, G is a (free) tree iff,


1. Any two vertices in G are connected by a unique simple path
2. G is connected, but if any edge is removed from E , the resulting graph is disconnected
3. G is connected, and $|E| = |V| - 1$
4. G is acyclic, and $|E| = |V| - 1$
5. G is acyclic, but if any edge is added to E , the resulting graph contains a cycle.

- In short, a connected acyclic undirected graph


Recall the data structure



Not tree
(graph)




Tree



Not tree
(forest)

Recall the data structure



- Rooted tree
 - A free tree in which one of the vertices is distinguished from the others (the root)
 - In a rooted tree, each node has **parent-child** and **ancestor-descendant** relationships.
 - The root has no parent/ancestor.

For today all keys are distinct.

Binary tree terminology

Each node has two **children**.

The **left child** of **3** is **2**

The **right child** of **3** is **4**

The **parent** of **3** is **5**

2 is a **descendant** of **5**


Each node has a pointer to its left child, right child, and parent.

Both **children** of **1** are NIL.
(I won't usually draw them).

The **height** of this tree is 3.
(Max number of edges from the root to a leaf).

This node is the **root**


This is a **node**.
It has a **key** (7).



These nodes are **leaves**.


Binary Search Trees

- It's a **binary search tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - (노드 x의 모든 왼쪽 후손들의 key는 x의 key 보다 작다)
 - Every RIGHT descendant of a node has key larger than that node.
 - (노드 x의 모든 오른쪽 후손들의 key는 x의 key 보다 크다)
- Example of building a binary search tree:



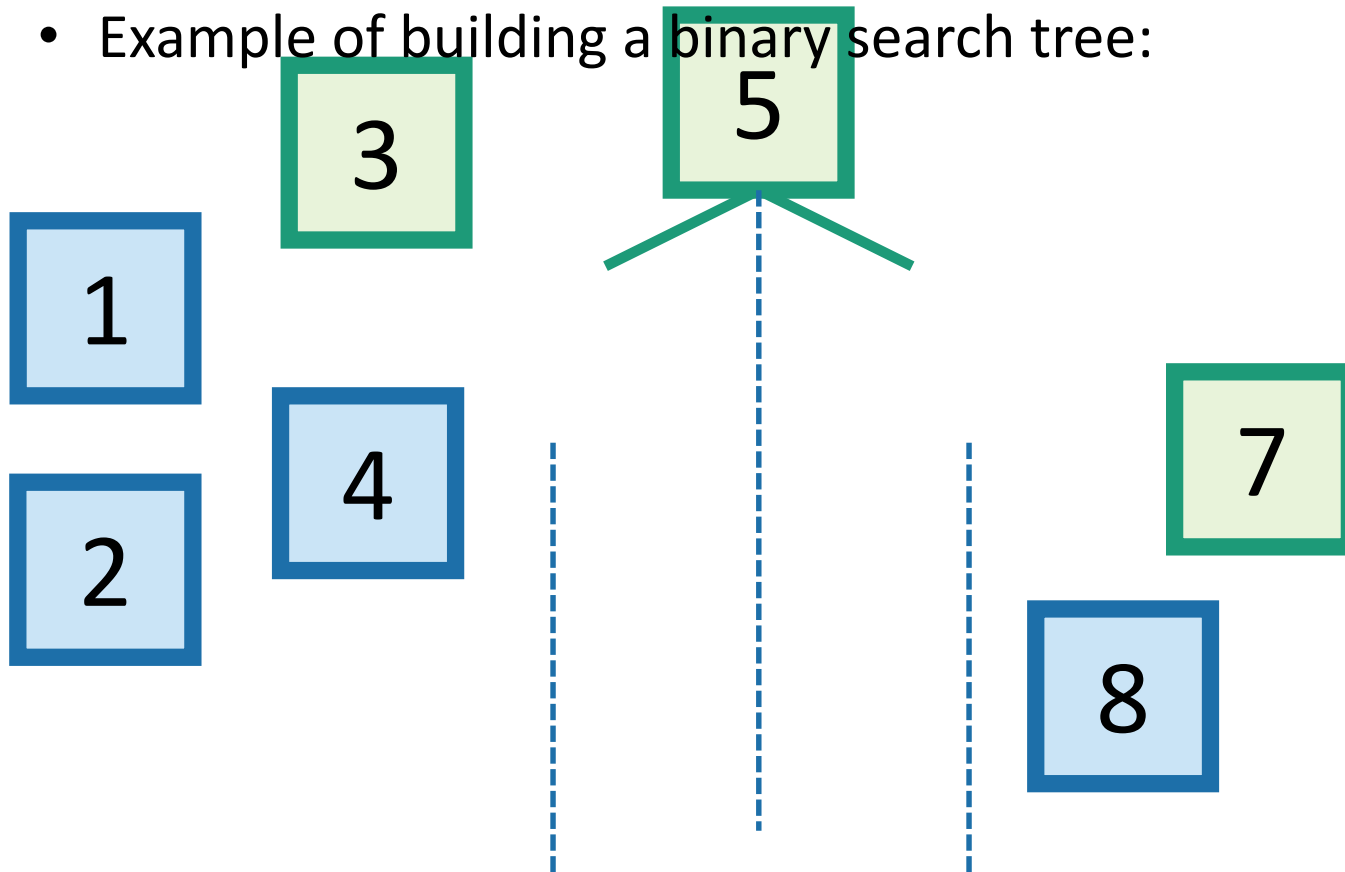
Binary Search Trees

- It's a **binary search tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - (노드 x의 모든 왼쪽 후손들의 key는 x의 key 보다 작다)
 - Every RIGHT descendant of a node has key larger than that node.
 - (노드 x의 모든 오른쪽 후손들의 key는 x의 key 보다 크다)
- Example of building a binary search tree:



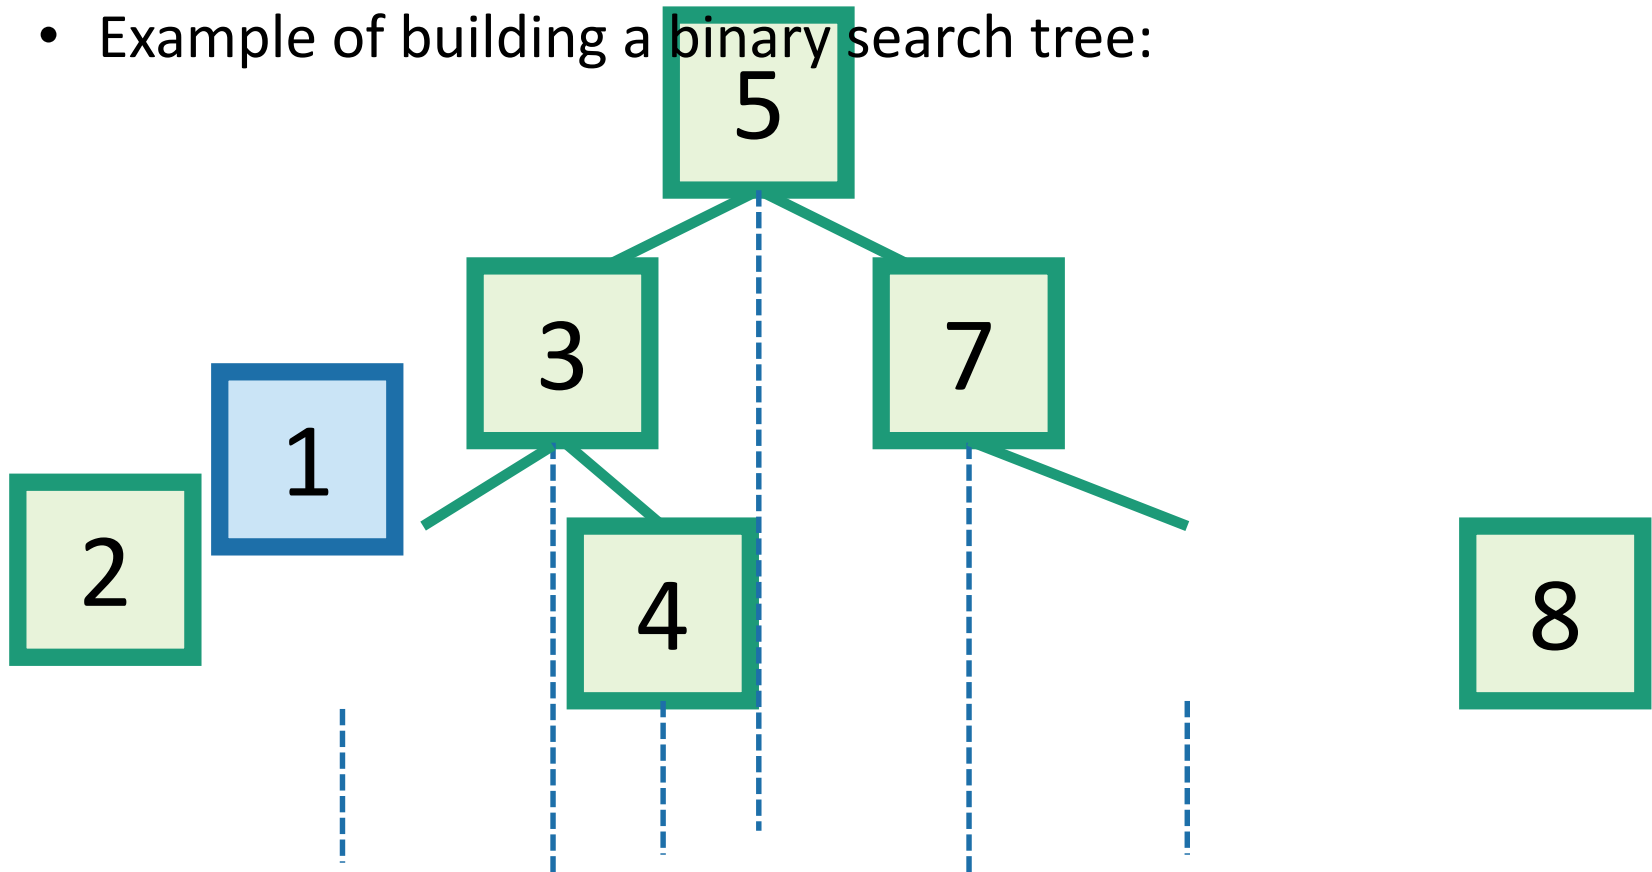
Binary Search Trees

- It's a **binary search tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - (노드 x의 모든 왼쪽 후손들의 key는 x의 key 보다 작다)
 - Every RIGHT descendant of a node has key larger than that node.
 - (노드 x의 모든 오른쪽 후손들의 key는 x의 key 보다 크다)
- Example of building a binary search tree:



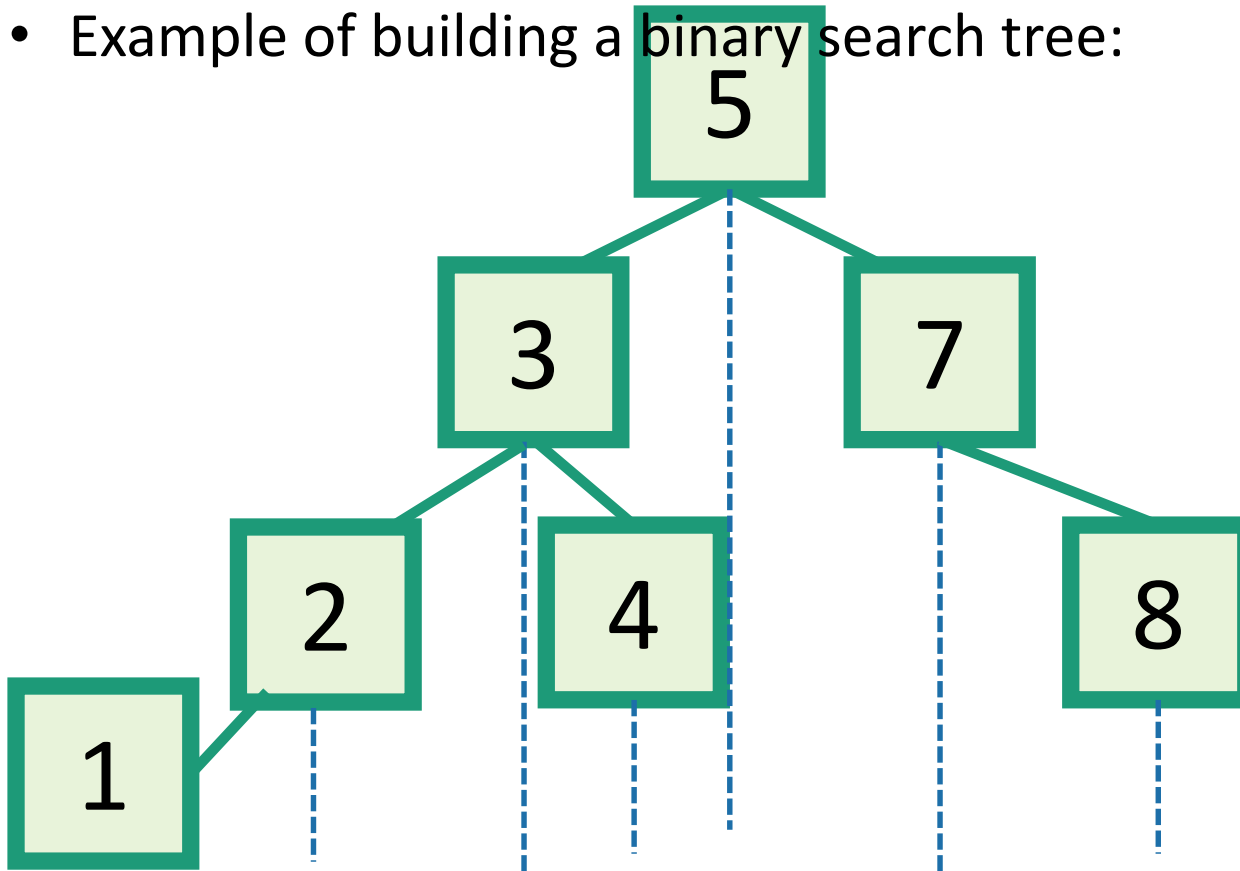
Binary Search Trees

- It's a **binary search tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - (노드 x의 모든 왼쪽 후손들의 key는 x의 key 보다 작다)
 - Every RIGHT descendant of a node has key larger than that node.
 - (노드 x의 모든 오른쪽 후손들의 key는 x의 key 보다 크다)
- Example of building a binary search tree:



Binary Search Trees

- It's a **binary search tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - (노드 x의 모든 왼쪽 후손들의 key는 x의 key 보다 작다)
 - Every RIGHT descendant of a node has key larger than that node.
 - (노드 x의 모든 오른쪽 후손들의 key는 x의 key 보다 크다)
- Example of building a binary search tree:




Q: Is this the only binary search tree I could possibly build with these values?

A: **No.** I made choices about which nodes to choose when. Any choices would have been fine.

Aside: this should look familiar



kinda like QuickSort



Binary Search Trees

- It's a **binary tree** so that:
 - **Every LEFT descendant** of a node has key less than that node.
 - **Every RIGHT descendant** of a node has key larger than that node.

Which of these is a BST?









Remember the goal

Fast **SEARCH/INSERT/DELETE**


Can we do these?

Motivation for Binary Search Trees

	Sorted Arrays	Linked Lists	Binary Search Trees*
Search	$O(\log(n))$ 	$O(n)$ 	$O(\log(n))$ 
Insert/Delete	$O(n)$ 	$O(1)$ 	$O(\log(n))$ 

SEARCH in a Binary Search Tree

definition by example



How long does this take?

$O(\text{length of longest path}) = O(\text{height})$


EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)

Write pseudocode
(or actual code) to
implement this!


INSERT in a Binary Search Tree



EXAMPLE: Insert 4.5

- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $\text{key} > x.\text{key}$:
 - Make a new node with the correct key, and put it as the right child of x .
 - **if** $\text{key} < x.\text{key}$:
 - Make a new node with the correct key, and put it as the left child of x .
 - **if** $x.\text{key} == \text{key}$:
 - **return**

DELETE in a Binary Search Tree



EXAMPLE: Delete 2


- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $x.\text{key} == \text{key}$:
 -delete x

This is a bit more complicated...

DELETE in a Binary Search Tree


several cases (by example)

say we want to delete 3



Case 1: if 3 is a leaf,
just delete it.

Write pseudocode for all of these!




Case 2: if 3 has just one child,
move that up.

DELETE in a Binary Search Tree

ctd.

Case 3: if 3 has two children,
replace 3 with its **immediate successor**.
(aka, next smallest thing after 3)



- Does this maintain the BST property?
 - Yes.
- How do we find the immediate successor?
 - SEARCH for 3 in the subtree under 3.right
- How do we remove it when we find it?
 - If [3.1] has 0 or 1 children, do one of the previous cases.
- What if [3.1] has two children?
 - It doesn't.

```
#include <iostream>
using namespace std;

struct node {
    int key;
    struct node *left, *right;
};

// Create a node
struct node *newNode(int item) {
    struct node *temp = (struct node *)malloc(sizeof(struct node));
    temp->key = item;
    temp->left = temp->right = NULL;
    return temp;
}
```

```
// Insert a node
struct node *insert(struct node *node, int key) {
    // Return a new node if the tree is empty
    if (node == NULL) return newNode(key);

    // Traverse to the right place and insert the node
    if (key < node->key)
        node->left = insert(node->left, key);
    else
        node->right = insert(node->right, key);

    return node;
}

struct node *search(struct node* node, int key) {
    if (node == NULL) return NULL;
    if (key == node->key) return node;
    else if (key < node->key) return search(node->left, key);
    else if (key > node->key) return search(node->right, key);
}
```

```

// Deleting a node
struct node *deleteNode(struct node *root, int key) {
    // Return if the tree is empty
    if (root == NULL) return root;

    // Find the node to be deleted
    if (key < root->key)
        root->left = deleteNode(root->left, key);
    else if (key > root->key)
        root->right = deleteNode(root->right, key);
    else {
        // If the node is with only one child or no child
        if (root->left == NULL) {
            struct node *temp = root->right;
            free(root);
            return temp;
        } else if (root->right == NULL) {
            struct node *temp = root->left;
            free(root);
            return temp;
        }

        // If the node has two children
        struct node *temp = minValueNode(root->right);

        // Place the inorder successor in position of the node to be deleted
        root->key = temp->key;

        // Delete the inorder successor
        root->right = deleteNode(root->right, temp->key);
    }
    return root;
}

```


```
// Find the inorder successor
struct node *minValueNode(struct node *node) {
    struct node *current = node;

    // Find the leftmost leaf
    while (current && current->left != NULL)
        current = current->left;

    return current;
}
```


How long do these operations take?

- **SEARCH** is the big one.
 - Everything else just calls **SEARCH** and then does some small $O(1)$ -time operation.



How long does search take?

Wait...



- This is a valid binary search tree.
- The version with n nodes has depth n , **not** $O(\log(n))$.

Could such a tree show up?
In what order would I have to
insert the nodes?

Inserting in the order
2,3,4,5,6,7,8 would do it.


So this **could** happen.

How often is “every so often” in the worst case?
It’s actually pretty often!

What to do?

- Goal: Fast **SEARCH/INSERT/DELETE**
- All these things take time $O(\text{height})$
- And the height might be big!!! ☹️
- Idea 0:
 - Keep track of how deep the tree is getting.
 - If it gets too tall, re-do everything from scratch.
 - At least $\Omega(n)$ every so often....
- Turns out that’s not a great idea. Instead we turn to...

Self-Balancing Binary Search Trees



Today's Outline

- Tree Algorithms
 - Tree basics, Binary search tree, Red-black tree
 - Reading: CLRS 12.1, 12.2, 12.3 and 13

How often is “every so often” in the worst case?
It’s actually pretty often!




What to do?

- Goal: Fast **SEARCH/INSERT/DELETE**
- All these things take time $O(\text{height})$
- And the height might be big!!! ☹️
- Idea 0:
 - Keep track of how deep the tree is getting.
 - If it gets too tall, re-do everything from scratch.
 - At least $\Omega(n)$ every so often....
- Turns out that’s not a great idea. Instead we turn to...



Idea 1: Rotations

No matter what lives underneath A,B,C,
this takes time $O(1)$. (Why?)


- Maintain Binary Search Tree (BST) property, while moving stuff around.





This seems helpful



Tree rotation



Tree rotation



Does this work?

- Whenever something seems unbalanced, do rotations until it's okay again.



Even for me this is pretty vague.
What do we mean by “seems unbalanced”?
What’s “okay”?

Idea 2: have some proxy for balance

- Maintaining **perfect balance** is too hard.
- Instead, come up with some **proxy for balance**:
 - If the tree satisfies **[SOME PROPERTY]**, then it's pretty balanced.
 - We can maintain **[SOME PROPERTY]** using rotations.




There are actually several ways to do this, but today we'll see...

Red-Black Trees

- A Binary Search Tree that balances itself!
- No more time-consuming by-hand balancing!

Red-Black tree!

Maintain balance by stipulating that **black nodes** are balanced, and that there *aren't too many red nodes*.




It's just good sense!

Red-Black Trees

these rules are the **proxy for balance**





- P1. Every node is colored **red** or **black**.
- P2. The root node is a **black node**.
- P3. NIL children count as **black nodes**.
- P4. Children of a **red node** are **black nodes**.
- P5. For all nodes x :
 - all paths from x to NIL's have the same number of **black nodes** on them.



I'm not going to draw the NIL children in the future, but they are treated as black nodes.


Examples(?)

- P1. Every node is colored **red** or **black**.
- P2. The root node is a **black node**.
- P3. NIL children count as **black nodes**.
- P4. Children of a **red node** are **black nodes**.
- P5. For all nodes x:
 - all paths from x to NIL's have the same number of **black nodes** on them.



Why??????

- This is **pretty balanced**.
 - The **black nodes** are balanced
 - The **red nodes** are “spread out” so they don’t mess things up too much.
- We can **maintain this property** as we insert/delete nodes, by using **rotations**.



This is the really clever idea!


This **Red-Black** structure is a **proxy for balance**.

It’s just a little weaker than perfect balance, but we can actually maintain it!



This is “pretty balanced”

- To see why, intuitively, let's try to build a Red-Black Tree that's **unbalanced**.




Guess: the height of a
red-black tree is at most $2 \log(n)$



That turns out to be basically right.

- Intuition:

- The height of completely balanced RBTree is $\text{floor}(\log n) + 1$.
- The number of black nodes from the root to any leaf is bounded to $\text{floor}(\log n) + 1$.
- And, the number of red nodes from the root to any leaf is less than $\text{floor}(\log n) + 1$.
- Then the longest path length is less than $2 * (\text{floor}(\log n) + 1)$.
- Time complexity of each operation: $O(\log n)$




Okay, so it's balanced...

...but can we maintain it?

- Yes!

- For the rest of lecture:
 - sketch of how we'd do this.
- See CLRS for more details.
- (You are not responsible for the details for this class – but you should understand the main ideas).


Many cases




- Suppose we want to insert **here**.
 - eg, want to insert 0.

Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.




What if it looks like this?



Example: insert 0




Many cases



- Suppose we want to insert **here**.
 - eg, want to insert 0.


Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.




What if it looks like this?

Example: insert 0




Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- **Fix things up if needed.**



What if it looks like this?




Example: insert 0

Can't we just insert 0 as a **black node**?




We need a bit more context




What if it looks like this?

Example: insert 0




We need a bit more context

- Add 0 as a red node.




What if it looks like this?

Example: insert 0



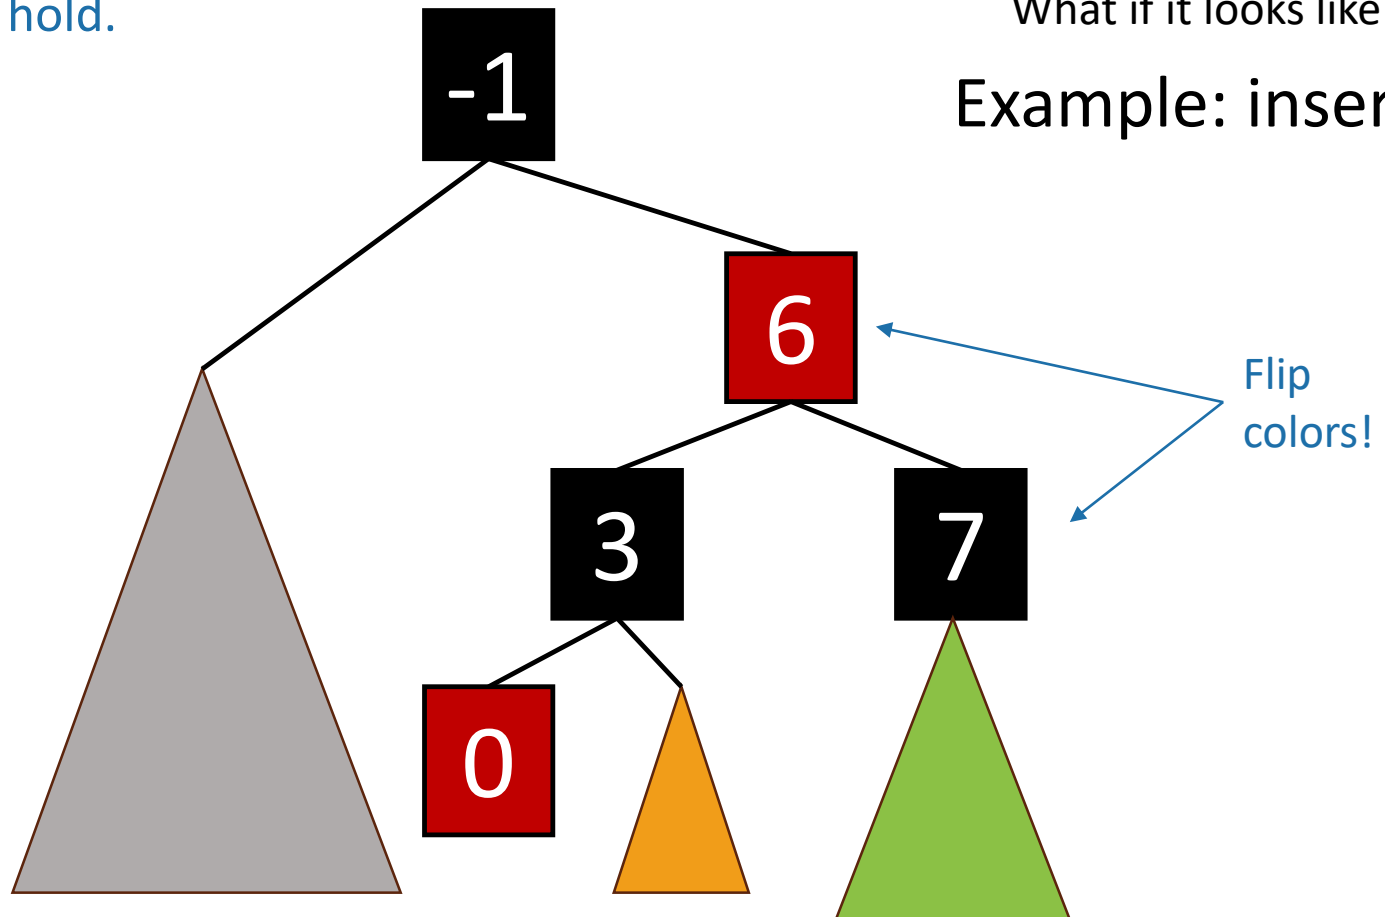
We need a bit more context

- Add 0 as a red node.
- **Claim:** RB-Tree properties still hold.






What if it looks like this?

Example: insert 0





But what if **that** was red?



What if it looks like this?

Example: insert 0



More context...



What if it looks like this?

Example: insert 0

More context...




What if it looks like this?

Example: insert 0


Now we're
basically inserting
6 into some
smaller tree.
Recurse!

Many cases




- Suppose we want to insert **here**.
 - eg, want to insert 0.

Inserting into a Red-Black Tree



What if it looks like this?

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.




Example: Insert 0.

- Actually, this can't happen?
 - 6-3 path has one black node
 - 6-7-... has at least two
- It might happen that we just turned 0 red from the previous step.
- Or it could happen if 7 is actually **NIL**.


Recall Rotations

- Maintain Binary Search Tree (BST) property, while moving stuff around.





Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.






What if it looks like this?

YOINK!



Need to argue that if RB-Tree property held before, it still does.

Many cases



- Suppose we want to insert **here**.
 - eg, want to insert 0.

insert in Red-Black Trees

```
enum Color {RED, BLACK};

struct Node
{
    int data;
    bool color;
    Node *left, *right, *parent;
};

void rotateLeft(Node *&root, Node *&pt)
{
    Node *pt_right = pt->right;
    pt->right = pt_right->left;

    if (pt->right != NULL)
        pt->right->parent = pt;

    pt_right->parent = pt->parent;
    if (pt->parent == NULL)
        root = pt_right;

    else if (pt == pt->parent->left)
        pt->parent->left = pt_right;

    else
        pt->parent->right = pt_right;

    pt_right->left = pt;
    pt->parent = pt_right;
}
```

insert in Red-Black Trees

```
// This function fixes violations caused by BST insertion
// It only contains a specific case. Other cases must be implemented!
void RBTree::fixViolation(Node *&root, Node *&pt)
{
    Node *parent_pt = NULL;
    Node *grand_parent_pt = NULL;


    while ((pt != root) && (pt->color != BLACK) &&
           (pt->parent->color == RED))
    {
        parent_pt = pt->parent;
        grand_parent_pt = pt->parent->parent;

        /* Case : 3
           Parent of pt is left child of Grand-parent of pt */
        if (parent_pt == grand_parent_pt->left)
        {
            Node *uncle_pt = grand_parent_pt->right;
            if (uncle_pt != NULL && uncle_pt->color == BLACK)
            {
                if (pt == parent_pt->left)
                {
                    rotateRight(root, grand_parent_pt);
                    swap(parent_pt->color, grand_parent_pt->color);
                    pt = parent_pt;
                }
            }
        }

        root->color = BLACK;
    }
}
```


Deleting from a Red-Black tree

- Some easy cases
 - If deleting node is RED
 - If deleting node is BLACK, and its unique child is RED



Deleting from a Red-Black tree

- Little-bit complicated
 - If deleting node is BLACK
 - In this case, we have address each case that (p,s,l,r) has distinct color combination (skip in this class)









That's a lot of cases

- You are **not responsible** for the details of Red-Black Trees. (For this class)
 - Though implementing them is a great exercise!
- You should know:
 - What are the properties of an RB tree?
 - And (more important) why does that guarantee that they are balanced?

What was the point again?


- Red-Black Trees **always** have height at most $2\log(n+1)$.
- As with general **Binary Search Trees**, all operations are $O(\text{height})$
- So all operations are $O(\log(n))$.

Conclusion: The best of both worlds

	Sorted Arrays	Linked Lists	Balanced Binary Search Trees
Search	$O(\log(n))$ 	$O(n)$ 	$O(\log(n))$ 
Insert/Delete	$O(n)$ 	$O(1)$ 	$O(\log(n))$ 

Today


- Begin a brief summary into **data structures!**



- Binary search trees
 - They are better when they're balanced.


this will lead us to...

- Self-Balancing Binary Search Trees
 - **Red-Black** trees.



Recap

- **Balanced binary trees** are the best of both worlds!
- But we need to **keep them balanced**.
- **Red-Black Trees** do that for us.
 - We get $O(\log(n))$ -time INSERT/DELETE/SEARCH
 - Clever idea: have a **proxy for balance**



**Post any question
On the Q&A board.**